

SUJAN DORA

AI Engineer

Contact

- United States
- +1 (940) 843-0811
- sujanworkofficial@gmail.com
- linkedin.com/in/sujan-dora
- github.com/sujandora
- My portfolio

TECHNICAL SKILLS

Programming Languages

- Python
- SQL

ML Frameworks & Libraries

- PyTorch
- vLLM
- Hugging Face

LLM & NLP Tools

- LLM Orchestration
- RAG
- Vector Search

Cloud & Deployment

- Distributed Processing
- Self-hosted Models

Databases & Data Engineering

- Vector Databases
- Entity Extraction

Web & DevOps

- API Integration
- Evaluation Frameworks

ML Concepts & Specializations

- Explainability
- Semantic Search
- Reasoning Workflows

Professional Summary

AI/ML Engineer with 2 years of hands-on experience delivering production-ready AI solutions, with a focus on LLM- and RAG-based document QA systems. Improved document QA accuracy by 20%, reduced MFA verification load by 25%, and streamlined ID authentication pipelines. Recent graduate with strong proficiency in Python, PyTorch, TensorFlow, OpenCV, and AWS, and a solid foundation in LLMs, RAG systems, model evaluation, and end-to-end ML workflows.

Education

Missouri State University Dec-2025
Master of Science in Computer Science — Data Science GPA: 3.63

Work Experience

AI/ML Engineer Jul 2023 – Dec 2023
Grootan Technologies Chennai, India

- Implemented a self-correcting RAG system to enable accurate, defensible insights from complex document sets, ensuring high explainability.
- Developed hybrid retrieval workflows (Dense + BM25) to improve search quality across large-scale heterogeneous corpora.
- Engineered batched query extraction pipelines, optimizing inference cost and latency for massive document analysis tasks.
- Built an LLM-as-Judge evaluation pipeline to benchmark model faithfulness and citation accuracy, critical for legal use cases.
- Designed structured extraction logic using Python to identify entities and issues within unstructured text headers.
- Deployed tailored inference infrastructure using Docker, supporting the integration of multiple model providers and self-hosted workflows.

AI/ML Intern Sep 2022 – Jul 2023
Grootan Technologies Chennai, India

- Built a National ID authentication system for detecting counterfeit holograms using YOLO object detection, MiDaS depth estimation, and GAN-generated synthetic samples to improve robustness.
- Optimized the hologram verification pipeline by eliminating redundant processing, reducing per-card detection time from 7 minutes to 3 minutes.
- Applied transfer learning to fine-tune pre-trained models for limited datasets, achieving reliable counterfeit detection across varied hologram designs.
- Containerized and deployed solutions using Docker and AWS, ensuring scalable production-ready pipelines for anomaly detection and identity verification.

Data Science Intern Nov 2021 – Mar 2022
iNeuron.ai Bengaluru, India

- Developed an end-to-end credit card default prediction project using Docker, CircleCI, and Heroku, achieving 92% model accuracy on simulated customer transactions.
- Built a CI/CD pipeline to deploy predictive models to AWS production, enabling real-time predictions.
- Conducted exploratory data analysis on 30,000+ customer records to identify key features that improved model performance.

Personal Projects

- BioMedical Research Assistant – Multi-Agent LLM System**
Fine-tuned Llama-3.1-8B on PubMedQA using QLoRA (4-bit) and built a CrewAI-based multi-agent biomedical QA pipeline with PubMed retrieval and a Streamlit interface.
- SafeSpace 2.0 – AI Mental Health Therapist**
Built an empathetic AI mental-health agent with crisis detection and therapist recommendations using LangGraph/REAct, deployed via Streamlit and Twilio WhatsApp with real-time API integrations.