SUJAN DORA

AI Engineer

Contact

San Jose, CA
+1 (940) 843-0811
sujanworkofficial@gmail.com
linkedin.com/in/sujan-dora
github.com/sujandora
My-portfolio

TECHNICAL SKILLS

Programming Languages
- Python
- SQL

ML Frameworks & Libraries
- TensorFlow
- PyTorch
- Hugging Face Transformers
- Scikit-learn

LLM & NLP Tools
- LangChain
- LangGraph
- LangSmith
- Multi-Agents
- LLM Agents
- RAG Systems
- OpenAI

Cloud & Deployment
- AWS
- Docker

Databases & AI Infrastructure
- PostgreSQL
- Vector Databases
- Weaviate

Web & DevOps
- FastAPI
- Git
- CI/CD Pipelines

ML Specializations
- Generative AI
- Agentic Workflows
- Deep Learning
- Hybrid RAG

Professional Summary

AI/ML Engineer with 2 years of hands-on experience delivering production-ready AI solutions, with a focus on LLM- and RAG-based document QA systems. Improved document QA accuracy by 20%, reduced MFA verification load by 25%, and streamlined ID authentication pipelines. Recent graduate with strong proficiency in Python, PyTorch, TensorFlow, OpenCV, and AWS, and a solid foundation in LLMs, RAG systems, model evaluation, and end-to-end ML workflows.

Education

Missouri State University — Dec-2025
Master of Science in Computer Science — Data Science — GPA: 3.63

Work Experience

AI/ML Engineer — Jul 2023 – Dec 2023
Grootan Technologies — Chennai, India
- Implemented a self-correcting RAG system using LangChain and Mistral 7B, improving answer reliability and mitigating context errors in multi-agent workflows.
- Built a hybrid retrieval workflow combining dense embeddings and BM25 search, enhancing relevance and coverage of retrieved documents for agentic reasoning.
- Developed an LLM-as-Judge evaluation pipeline to assess faithfulness and relevance across hundreds of QA examples, ensuring reliable performance of autonomous agents.
- Introduced batched query extraction using Python and SQL, reducing LLM calls by 80
- Designed an ML-based anomaly detection system for workflows, reducing redundant prompts by 25
- Containerized and deployed scalable AI agent solutions using Docker and AWS, ensuring robust pipelines for collaborative agent systems.

AI/ML Intern — Sep 2022 – Jul 2023
Grootan Technologies — Chennai, India
- Built a National ID authentication system for detecting counterfeit holograms using YOLO object detection, MiDaS depth estimation, and GAN-generated synthetic samples to improve robustness.
- Optimized the hologram verification pipeline by eliminating redundant processing, reducing per-card detection time from 7 minutes to 3 minutes.
- Applied transfer learning to fine-tune pre-trained models for limited datasets, achieving reliable counterfeit detection across varied hologram designs.
- Containerized and deployed solutions using Docker and AWS, ensuring scalable production-ready pipelines for anomaly detection and identity verification.

Data Science Intern — Nov 2021 – Mar 2022
iNeuron.ai — Bengaluru, India
- Developed an end-to-end credit card default prediction project using Docker, CircleCI, and Heroku, achieving 92% model accuracy on simulated customer transactions.
- Built a CI/CD pipeline to deploy predictive models to AWS production, enabling real-time predictions.
- Conducted exploratory data analysis on 30,000+ customer records to identify key features that improved model performance.

Personal Projects

- **BioMedical Research Assistant – Multi-Agent LLM System**
  Fine-tuned Llama-3.1-8B on PubMedQA using QLoRA (4-bit) and built a CrewAI-based multi-agent biomedical QA pipeline with PubMed retrieval and a Streamlit interface.
- **SafeSpace 2.0 – AI Mental Health Therapist**
  Built an empathetic AI mental-health agent with crisis detection and therapist recommendations using LangGraph/REAct, deployed via Streamlit and Twilio WhatsApp with real-time API integrations.