

SUJAN DORA

AI Engineer

Contact

- 📍 San Francisco, CA (Open to relocate)
- ☎ +1 (940) 843-0811
- ✉ sujanworkofficial@gmail.com
- 🌐 linkedin.com/in/sujan-dora
- 🐙 github.com/sujandora
- 📁 My-portfolio

TECHNICAL SKILLS

Programming Languages

- Python
- SQL

ML Frameworks & Libraries

- TensorFlow
- PyTorch
- Scikit-learn
- Hugging Face Transformers

LLM & NLP Tools

- LangChain
- LangSmith
- LlamaIndex
- Multi-Agents
- Finetuning (LoRA, QLoRA)
- OpenAI, Ollama
- RAG Systems
- Multimodal AI

ML Specializations

- Deep Learning
- Computer Vision
- Anomaly Detection

MLOps & Deployment

- MLflow
- Git, GitHub Actions
- CI/CD Pipelines
- Model Deployment

Cloud & Infrastructure

- AWS
- Azure

Databases & AI Infrastructure

- PostgreSQL, MySQL, MongoDB
- Weaviate
- FAISS
- ChromaDB

Web & DevOps

- FastAPI
- Docker Containerization

Professional Summary

AI/ML Engineer with 2+ years of hands-on experience delivering production-ready AI solutions, with a focus on LLM- and RAG-based document QA systems. Improved document QA accuracy by 20%, reduced MFA verification load by 25%, and streamlined ID authentication pipelines. Recent graduate with strong proficiency in Python, PyTorch, TensorFlow, OpenCV, and AWS, and a solid foundation in LLMs, RAG systems, model evaluation, and end-to-end ML workflows.

Education

Missouri State University Dec-2025
Master of Science in Computer Science — Data Science GPA: 3.63

Work Experience

AI/ML Engineer Jul 2023 – Dec 2023
Grootan Technologies Chennai, India

- Architected and deployed a self-correcting RAG system using LangChain and Mistral 7B for enterprise-scale GenAI applications, improving document QA accuracy by ~20% and enabling advanced AI agents with multimodal functionality and tool-calling capabilities.
- Engineered hybrid retrieval architecture combining dense embeddings (FAISS) with BM25 search for frontier data applications, achieving 35% improvement in retrieval relevance for enterprise clients across cybersecurity, journalism, and genomics domains.
- Developed LLM-as-Judge evaluation framework to assess faithfulness, performance, and business impact across 500+ GenAI use cases, ensuring production-ready reliability and data-driven optimization for next-generation AI experiences.
- Optimized enterprise LLM pipelines through batched query extraction and tool-calling agents, reducing API costs by ~80% and cutting inference latency by 60% while maintaining scalability for transformative client applications.
- Designed ML-based anomaly detection system for AI security firewalls, reducing false positives by 25% while maintaining threat detection coverage and improving operational efficiency in production AI deployments.
- Built a multi-agent research workflow using CrewAI with proprietary Scale platform integration, enabling complex AI agent development with autonomous decision-making and cross-functional collaboration for enterprise customer solutions.

AI/ML Intern Jun 2022 – Jul 2023
Grootan Technologies Chennai, India

- Built a National ID authentication system for detecting counterfeit holograms using YOLO object detection, MiDaS depth estimation, and GAN-generated synthetic samples to improve robustness.
- Optimized the hologram verification pipeline by eliminating redundant processing, reducing per-card detection time from 7 minutes to 3 minutes.
- Applied transfer learning to fine-tune pre-trained models for limited datasets, achieving reliable counterfeit detection across varied hologram designs.
- Containerized and deployed solutions using Docker and AWS, ensuring scalable production-ready pipelines for anomaly detection and identity verification.

Data Science Intern Nov 2021 – May 2022
iNeuron.ai Bengaluru, India

- Developed an end-to-end credit card default prediction project using Docker, CircleCI, and Heroku, achieving 92% model accuracy on simulated customer transactions.
- Built a CI/CD pipeline to deploy predictive models to AWS production, enabling real-time predictions.
- Conducted exploratory data analysis on 30,000+ customer records to identify key features that improved model performance.

Personal Projects

- **BioMedical Research Assistant – Multi-Agent LLM System**
Fine-tuned Llama-3.1-8B on PubMedQA using QLoRA (4-bit) and built a CrewAI-based multi-agent biomedical QA pipeline with PubMed retrieval and a Streamlit interface.
- **SafeSpace 2.0 – AI Mental Health Therapist**
Built an empathetic AI mental-health agent with crisis detection and therapist recommendations using LangGraph/REAct, deployed via Streamlit and Twilio WhatsApp with

real-time API integrations.